

5133178225132135225133181225134171

May 31, 2024

We need to install libraries below.

```
[50]: !pip install transformers datasets evaluate sequeval tqdm
```

Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages (4.41.1)

Requirement already satisfied: datasets in /usr/local/lib/python3.10/dist-packages (2.19.1)

Requirement already satisfied: evaluate in /usr/local/lib/python3.10/dist-packages (0.4.2)

Requirement already satisfied: sequeval in /usr/local/lib/python3.10/dist-packages (1.2.2)

Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (4.66.4)

Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers) (3.14.0)

Requirement already satisfied: huggingface-hub<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.23.1)

Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (1.25.2)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (24.0)

Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (6.0.1)

Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (2024.5.15)

Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from transformers) (2.31.0)

Requirement already satisfied: tokenizers<0.20,>=0.19 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.19.1)

Requirement already satisfied: safetensors>=0.4.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.4.3)

Requirement already satisfied: pyarrow>=12.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (14.0.2)

Requirement already satisfied: pyarrow-hotfix in /usr/local/lib/python3.10/dist-packages (from datasets) (0.6)

Requirement already satisfied: dill<0.3.9,>=0.3.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.3.8)

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from datasets) (2.0.3)

Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages (from datasets) (3.4.1)

Requirement already satisfied: multiprocessing in /usr/local/lib/python3.10/dist-packages (from datasets) (0.70.16)

Requirement already satisfied: fsspec[http]<=2024.3.1,>=2023.1.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (2023.6.0)

Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets) (3.9.5)

Requirement already satisfied: scikit-learn>=0.21.3 in /usr/local/lib/python3.10/dist-packages (from sequeval) (1.2.2)

Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.3.1)

Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (23.2.0)

Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.4.1)

Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (6.0.5)

Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.9.4)

Requirement already satisfied: async-timeout<5.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (4.0.3)

Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0,>=0.23.0->transformers) (4.11.0)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (3.7)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (2.0.7)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->transformers) (2024.2.2)

Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.21.3->sequeval) (1.11.4)

Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.21.3->sequeval) (1.4.2)

Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.21.3->sequeval) (3.5.0)

Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2023.4)

Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2024.1)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->datasets) (1.16.0)

## 0. CONFIGURATION

```
[74]: class DotDict(dict):
        """dot notation access to dictionary attributes"""
        __getattr__ = dict.get
        __setattr__ = dict.__setitem__
        __delattr__ = dict.__delitem__

config = {
##### EDIT
↪#####
    'seed' : 2024,
    'batch_size' : 32,
    'lr' : 0.0001,
    'weight_decay' : 0.01,
    'hidden_size' : 512,
    'num_heads' : 8,
    'num_encoder_layers' : 6,
    'hidden_dropout_prob' : 0.1,
    'use_lstm': False,
##### EDIT
↪#####
    'num_epochs' : 5, # NEVER TOUCH
    'vocab_size' : 30522, # NEVER TOUCH
    'pad_token_id' : 0, # NEVER TOUCH
    'num_labels' : 9, # NEVER TOUCH
}

config = DotDict(config)
```

```
[52]: import torch
import random
import warnings
warnings.filterwarnings('ignore')

import numpy as np

def set_seed(config):
    random.seed(config.seed)
    np.random.seed(config.seed)
    torch.manual_seed(config.seed)
    torch.cuda.manual_seed(config.seed)
    torch.cuda.manual_seed_all(config.seed)
    torch.backends.cudnn.benchmark = True
```

```
torch.backends.cudnn.deterministic = True

set_seed(config)
```

## 1. DATASET

```
[53]: from datasets import load_dataset

dataset = load_dataset("conll2003")
# dataset
```

The letter that prefixes each ner\_tag indicates the token position of the entity:

B- indicates the beginning of an entity.

I- indicates a token is contained inside the same entity (for example, the State token is a part of an entity like Empire State Building).

O indicates the token doesn't correspond to any entity.

```
[54]: label_list = dataset['train'].features['ner_tags'].feature.names
config.num_labels = len(label_list)
# label_list
```

```
[55]: from transformers import AutoTokenizer

tokenizer = AutoTokenizer.from_pretrained('distilbert/distilbert-base-uncased')
config.vocab_size = len(tokenizer.get_vocab())
config.pad_token_id = tokenizer.get_vocab()['[PAD]']
```

```
[56]: example = dataset["test"][0]
# example
```

```
[57]: tokenized_input = tokenizer(example["tokens"], is_split_into_words=True)
# tokenized_input
```

```
[58]: tokens = tokenizer.convert_ids_to_tokens(tokenized_input["input_ids"])
# tokens
```

```
[87]: def tokenize_and_align_labels(examples):
    tokenized_inputs = tokenizer(examples["tokens"], truncation=True,
    ↪is_split_into_words=True)
    labels = []

    for i, label in enumerate(examples[f"ner_tags"]):
        # Map tokens to their respective word.
        # Set the special tokens to -100.
        # Only label the first token of a given word.
```

```

##### EDIT
#####

word_ids = tokenized_inputs.word_ids(batch_index=i)
label_ids = []
previous_word_idx = None
for word_idx in word_ids:
    if word_idx is None:
        label_ids.append(-100)
    elif word_idx != previous_word_idx:
        label_ids.append(label[word_idx])
    else:
        label_ids.append(-100)
    previous_word_idx = word_idx
labels.append(label_ids)

##### EDIT
#####

tokenized_inputs["labels"] = labels
return tokenized_inputs

```

```

[76]: tokenized_dataset = dataset.map(tokenize_and_align_labels, batched=True,
    ↪ remove_columns=dataset["train"].column_names)

```

```

Map:   0%|          | 0/14041 [00:00<?, ? examples/s]
Map:   0%|          | 0/3250 [00:00<?, ? examples/s]
Map:   0%|          | 0/3453 [00:00<?, ? examples/s]

```

```

[77]: from transformers import DataCollatorForTokenClassification

data_collator = DataCollatorForTokenClassification(tokenizer=tokenizer)

```

```

[78]: from torch.utils.data import DataLoader

train_dataloader = DataLoader(
    tokenized_dataset["train"],
    shuffle=True,
    collate_fn=data_collator,
    batch_size=config.batch_size,
)

valid_dataloader = DataLoader(
    tokenized_dataset["validation"],
    shuffle=True,
    collate_fn=data_collator,
)

```

```

        batch_size=config.batch_size,
    )

    test_dataloader = DataLoader(
        tokenized_dataset["test"],
        shuffle=False,
        collate_fn=data_collator,
        batch_size=config.batch_size,
    )

```

## 2. METRIC

```
[79]: import evaluate
```

```
sequeval = evaluate.load("sequeval")
```

```
[80]: def compute_metrics(all_predictions, all_labels):
    predictions, labels = None, None

    # For evaluation, postprocess for shape.
    ##### EDIT
    #####

    predictions, labels = [], []

    for batch_preds, batch_labels in zip(all_predictions, all_labels):
        for preds, labs in zip(batch_preds, batch_labels):
            predictions.append(preds)
            labels.append(labs)

    ##### EDIT
    #####

    true_predictions = [
        [label_list[p] for (p, l) in zip(prediction, label) if l != -100]
        for prediction, label in zip(predictions, labels)
    ]

    true_labels = [
        [label_list[l] for (p, l) in zip(prediction, label) if l != -100]
        for prediction, label in zip(predictions, labels)
    ]

    results = sequeval.compute(predictions=true_predictions,
    references=true_labels)
    return {
        "precision": results["overall_precision"],

```

```

    "recall": results["overall_recall"],
    "f1": results["overall_f1"],
    "accuracy": results["overall_accuracy"],
}

```

### 3. MODEL

```

[81]: from torch import nn

from typing import Optional

def _expand_mask(mask, tgt_len = None):
    """
    Inputs
        mask.shape = (B, S_L)
    Outputs
        output.shape = (B, 1, T_L, S_L)
    """
    batch_size, src_len = mask.size()
    tgt_len = tgt_len if tgt_len is not None else src_len

    expanded_mask = mask[:, None, None, :].expand(batch_size, 1, tgt_len,
↪src_len).to(torch.float)

    inverted_mask = 1.0 - expanded_mask

    return inverted_mask.masked_fill(inverted_mask.bool(), torch.finfo(torch.
↪float).min)

class PositionalEncoding(nn.Module):
    def __init__(self, hidden_size):
        super().__init__()
        self.encoding = torch.zeros(1024, hidden_size)
        self.encoding.requires_grad = False

        pos = torch.arange(0, 1024)
        pos = pos.float().unsqueeze(dim = 1)

        _2i = torch.arange(0, hidden_size, step = 2).float()

        ##### EDIT ↵
↪#####

        #         positional encoding
        self.encoding[:, 0::2] = torch.sin(pos / (10000 ** (_2i / hidden_size)))

```

```

        self.encoding[:, 1::2] = torch.cos(pos / (10000 ** (_2i / hidden_size)))

        ##### EDIT
    <--#####

    def forward(self, x):
        batch_size, seq_len = x.size()
        device = x.device

        return self.encoding[:seq_len, :].unsqueeze(0).to(device)

class MultiHeadAttention(nn.Module):
    def __init__(self, config):
        super().__init__()
        self.hidden_size = config.hidden_size
        self.num_heads = config.num_heads
        self.d_head = self.hidden_size // self.num_heads
        self.scaling = self.d_head ** -0.5

        self.q_proj = nn.Linear(self.hidden_size, self.hidden_size)
        self.k_proj = nn.Linear(self.hidden_size, self.hidden_size)
        self.v_proj = nn.Linear(self.hidden_size, self.hidden_size)
        self.out_proj = nn.Linear(self.hidden_size, self.hidden_size)
        self.dropout = nn.Dropout(0.1)

    def _shape(self, tensor, seq_len, batch_size):
        return tensor.view(batch_size, seq_len, self.num_heads, self.d_head).
    <--transpose(1, 2).contiguous()

    def forward(self, query_states, key_value_states, attention_mask):
        attn_output = None

        ##### EDIT
    <--#####

        # Query, key, value projection
        query_states = self.q_proj(query_states) * self.scaling
        key_states = self.k_proj(key_value_states)
        value_states = self.v_proj(key_value_states)

        # Multi-head
        batch_size, seq_len, _ = query_states.size()
        query_states = self._shape(query_states, seq_len, batch_size)
        key_states = self._shape(key_states, seq_len, batch_size)
        value_states = self._shape(value_states, seq_len, batch_size)

```



```

        # Attention weight
        attn_weights = torch.matmul(query_states, key_states.transpose(-1, -2))
        expanded_attention_mask = _expand_mask(attention_mask, seq_len) #
        ↪
        attn_weights = attn_weights + expanded_attention_mask
        attn_weights = torch.nn.functional.softmax(attn_weights, dim=-1)
        attn_weights = self.dropout(attn_weights)

        # Attention
        attn_output = torch.matmul(attn_weights, value_states)

        # Multi-head
        attn_output = attn_output.transpose(1, 2).contiguous().view(batch_size, ↪
        ↪seq_len, self.hidden_size)

        # Output projection
        attn_output = self.out_proj(attn_output)

        ##### EDIT ↪
        ↪#####

        return attn_output

class EncoderLayer(nn.Module):
    def __init__(self, config):
        super().__init__()
        self.hidden_size = config.hidden_size

        self.self_attn = MultiHeadAttention(config)
        self.self_attn_layer_norm = nn.LayerNorm(self.hidden_size)

        self.activation_fn = nn.ReLU()

        self.fc1 = nn.Linear(self.hidden_size, 4 * self.hidden_size)
        self.fc2 = nn.Linear(4 * self.hidden_size, self.hidden_size)
        self.final_layer_norm = nn.LayerNorm(self.hidden_size)

        self.dropout = nn.Dropout(0.1)

    def forward(self, hidden_states, enc_self_mask):
        residual = hidden_states
        hidden_states = self.self_attn(
            query_states = hidden_states,
            key_value_states = hidden_states,
            attention_mask = enc_self_mask
        )

```

```

        hidden_states = self.dropout(hidden_states)
        hidden_states = residual + hidden_states
        hidden_states = self.self_attn_layer_norm(hidden_states)

        residual = hidden_states
        hidden_states = self.activation_fn(self.fc1(hidden_states))
        hidden_states = self.dropout(hidden_states)
        hidden_states = self.fc2(hidden_states)
        hidden_states = self.dropout(hidden_states)
        hidden_states = residual + hidden_states
        hidden_states = self.final_layer_norm(hidden_states)

    return hidden_states

class Encoder(nn.Module):
    def __init__(self, config, embed_tokens, embed_positions):
        super().__init__()
        self.hidden_size = config.hidden_size

        self.embed_tokens = embed_tokens
        self.embed_positions = embed_positions

        self.layers = nn.ModuleList([EncoderLayer(config) for _ in range(config.
↪ num_encoder_layers)])
        self.embedding_layer_norm = nn.LayerNorm(self.hidden_size)

    def forward(self, enc_ids, enc_mask):
        enc_hidden_states = None

        ##### EDIT_
↪ #####

        # ID positional encoding
        enc_hidden_states = self.embed_tokens(enc_ids) + self.
↪ embed_positions(enc_ids)
        enc_hidden_states = self.embedding_layer_norm(enc_hidden_states)

        # Encoder layer
        for layer in self.layers:
            enc_hidden_states = layer(enc_hidden_states, enc_mask)

        ##### EDIT_
↪ #####

    return enc_hidden_states

```

```

class TransformerEncoder(nn.Module):
    def __init__(self, config):
        super().__init__()
        self.embed_tokens = nn.Embedding(config.vocab_size, config.hidden_size,
        ↪config.pad_token_id)
        self.embed_positions = PositionalEncoding(config.hidden_size)
        self.encoder = Encoder(config, self.embed_tokens, self.embed_positions)

    def forward(self, enc_ids, enc_mask = None):
        enc_hidden_states = self.encoder(enc_ids, enc_mask)

        return enc_hidden_states

class LSTMCell(nn.Module):
    def __init__(self, config):
        super().__init__()
        self.input_size = config.hidden_size
        self.hidden_size = config.hidden_size

        self.linear_ih = nn.Linear(self.input_size, 4 * self.hidden_size)
        self.linear_hh = nn.Linear(self.hidden_size, 4 * self.hidden_size)

    def forward(self, input, hidden):
        hy, cy = None, None

        ##### EDIT ↪
        ↪#####

        hx, cx = hidden

        # input hidden state linear transformation
        gates = self.linear_ih(input) + self.linear_hh(hx)

        # gate input
        ingate, forgetgate, cellgate, outgate = gates.chunk(4, 1)

        #
        ingate = torch.sigmoid(ingate)
        forgetgate = torch.sigmoid(forgetgate)
        cellgate = torch.tanh(cellgate)
        outgate = torch.sigmoid(outgate)

        # cell state, hidden state
        cy = (forgetgate * cx) + (ingate * cellgate)
        hy = outgate * torch.tanh(cy)

```

```

##### EDIT
#####

    return hy, cy

class ModelForNER(nn.Module):
    def __init__(self, config):
        super().__init__()
        self.transformer_encoder = TransformerEncoder(config)
        self.lstm = LSTMCell(config)

        self.dropout = nn.Dropout(config.hidden_dropout_prob)
        self.classifier = nn.Linear(config.hidden_size, config.num_labels)
        self.num_labels = config.num_labels
        self.use_lstm = config.use_lstm

    def forward(self,
                input_ids: Optional[torch.Tensor] = None,
                attention_mask: Optional[torch.Tensor] = None,
                ):

        logits = None

        ##### EDIT
        #####

        if self.use_lstm:
            # LSTM hidden states
            enc_hidden_states = self.transformer_encoder(input_ids,
            attention_mask)
            batch_size = enc_hidden_states.size(0)
            hx, cx = torch.zeros(batch_size, self.hidden_size,
            device=enc_hidden_states.device), torch.zeros(batch_size, self.hidden_size,
            device=enc_hidden_states.device)
            outputs = []
            for t in range(enc_hidden_states.size(1)):
                hx, cx = self.lstm(enc_hidden_states[:, t, :], (hx, cx))
                outputs.append(hx)
            # LSTM classifier
            logits = self.classifier(torch.stack(outputs, dim=1))
        else:
            # Transformer Encoder classifier
            enc_hidden_states = self.transformer_encoder(input_ids,
            attention_mask)

```

```

        logits = self.classifier(enc_hidden_states)

        ##### EDIT
        #####

    return logits

```

```
[82]: model = ModelForNER(config=config)
```

#### 4. Optimizer & Scheduler

```
[83]: from torch.optim import AdamW
      from torch.nn import CrossEntropyLoss

      optimizer = AdamW(model.parameters(), lr=config.lr,
                        weight_decay=config.weight_decay)
      criterion = CrossEntropyLoss()
```

```
[84]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
      model = model.to(device)
```

#### 5. TRAINING

```
[85]: from tqdm.auto import tqdm

      for epoch in range(config.num_epochs):
          print(f'##### EPOCH: {epoch} #####')
          model.train()

          # TRAINING LOOP
          for i, data in tqdm(enumerate(train_dataloader),
                              total=len(train_dataloader),
                              desc='TRAINING'):

              input_ids = data['input_ids'].to(device)
              attention_mask = data['attention_mask'].to(device)
              labels = data['labels'].to(device)

              optimizer.zero_grad()
              outputs = model(input_ids, attention_mask)
              loss = criterion(outputs.view(-1, outputs.shape[-1]), labels.view(-1))

              loss.backward()
              optimizer.step()

          # VALIDATION LOOP
          model.eval()
```

```

valid_loss = 0.0
all_predictions, all_labels = [], []

with torch.no_grad():
    for i, data in tqdm(enumerate(valid_dataloader),
        ↪total=len(valid_dataloader), desc='VALIDATION'):
        input_ids = data['input_ids'].to(device)
        attention_mask = data['attention_mask'].to(device)
        labels = data['labels'].to(device)

        outputs = model(input_ids, attention_mask)
        loss = criterion(outputs.view(-1, outputs.shape[-1]), labels.
        ↪view(-1))
        valid_loss += loss.item()

        all_predictions.append(np.argmax(outputs.cpu().numpy(), axis=2))
        all_labels.append(labels.cpu().numpy())

    print(f'Validation Loss: {valid_loss / len(valid_dataloader)}')

    metrics = compute_metrics(all_predictions, all_labels)
    print(f'Validation metrics: {metrics}')

print('Finished Training')

```

```

##### EPOCH: 0 #####
TRAINING:   0%|          | 0/439 [00:00<?, ?it/s]
VALIDATION: 0%|          | 0/102 [00:00<?, ?it/s]

Validation Loss: 0.3491189303059204
Validation metrics: {'precision': 0.5636295180722891, 'recall':
0.5038707505890273, 'f1': 0.5320774835614004, 'accuracy': 0.9059616058564698}
##### EPOCH: 1 #####
TRAINING:   0%|          | 0/439 [00:00<?, ?it/s]
VALIDATION: 0%|          | 0/102 [00:00<?, ?it/s]

Validation Loss: 0.2792895826346734
Validation metrics: {'precision': 0.5989261308167914, 'recall':
0.6194883877482329, 'f1': 0.6090337524818, 'accuracy': 0.9221798216580351}
##### EPOCH: 2 #####
TRAINING:   0%|          | 0/439 [00:00<?, ?it/s]
VALIDATION: 0%|          | 0/102 [00:00<?, ?it/s]

Validation Loss: 0.27210846397222255
Validation metrics: {'precision': 0.6267673521850899, 'recall':

```

```

0.6565129585997981, 'f1': 0.6412954134473121, 'accuracy': 0.9294030606284802}
##### EPOCH: 3 #####
TRAINING:  0%|          | 0/439 [00:00<?, ?it/s]
VALIDATION: 0%|          | 0/102 [00:00<?, ?it/s]
Validation Loss: 0.27487604561097484
Validation metrics: {'precision': 0.6423742786479802, 'recall':
0.6556714910804443, 'f1': 0.6489547763804447, 'accuracy': 0.9342704723336319}
##### EPOCH: 4 #####
TRAINING:  0%|          | 0/439 [00:00<?, ?it/s]
VALIDATION: 0%|          | 0/102 [00:00<?, ?it/s]
Validation Loss: 0.2769594073441683
Validation metrics: {'precision': 0.6664976335361731, 'recall':
0.6635812857623695, 'f1': 0.6650362624388598, 'accuracy': 0.9346014563295821}
Finished Training

```

```

[86]: # TEST
model.eval()
all_predictions, all_labels = [], []

with torch.no_grad():
    for i, data in tqdm(enumerate(test_dataloader), total=len(test_dataloader),
        desc='TEST'):
        input_ids = data['input_ids'].to(device)
        attention_mask = data['attention_mask'].to(device)
        labels = data['labels'].to(device)

        outputs = model(input_ids, attention_mask)
        loss = criterion(outputs.view(-1, outputs.shape[-1]), labels.view(-1))

        all_predictions.append(np.argmax(outputs.cpu().numpy(), axis=2))
        all_labels.append(labels.cpu().numpy())

metrics = compute_metrics(all_predictions, all_labels)
print(f'Test metrics: {metrics}')

```

```

TEST:  0%|          | 0/108 [00:00<?, ?it/s]
Test metrics: {'precision': 0.5969534050179212, 'recall': 0.5897662889518414,
'f1': 0.593338083363021, 'accuracy': 0.9179282868525896}

```